



Import notebook

Import Libraries

```
# Pandas is a Python data analysis library used for analysing data
import pandas as pd

# Is a Python library used for numerical analysis/to perform mathematical calculations
import numpy as np
```

Data Ingestion

```
data_path = '/Workspace/Users/michaelarobyn01@gmail.com/Patient Readmission Analysis/Expanded_Patient_Readmission_Data (1).csv'
```

```
# To read the data from a csv file to a pandas DataFrame
df=pd.read_csv(data_path)
```

```
> df: pandas.core.frame.DataFrame = [Patient ID: int64, Age: int64 ... 8 more fields]
```

```
# Pandas function to help us to see the first 10 rows of the data
df.head(10)
```

	Patient ID	Age	Gender	Admission Type	Length of Stay	Number of Diagnoses	Blood Pressure	Blood Sugar Levels	Previous Admissions	Readmission
0	1	62	Female	Elective	4	5	110	130	1	No
1	2	65	Male	Emergency	19	2	157	81	4	No
2	3	82	Female	Emergency	18	4	74	84	0	No
3	4	85	Male	Emergency	2	4	106	85	4	No
4	5	85	Female	Elective	19	3	80	119	3	No
5	6	27	Male	Emergency	18	6	136	99	2	No
6	7	39	Male	Elective	21	3	116	74	0	Yes
7	8	54	Male	Emergency	3	4	136	68	3	No
8	9	88	Female	Elective	6	6	120	81	2	No
9	10	30	Male	Elective	6	8	139	106	3	No

EDA

Step 1: Understand the data structure

```
# To check number of rows and columns of the data
df.shape
```

```
(3000, 10)
```

```
df.columns
Index(['Patient ID', 'Age', 'Gender', 'Admission Type', 'Length of Stay',
       'Number of Diagnoses', 'Blood Pressure', 'Blood Sugar Levels',
       'Previous Admissions', 'Readmission'],
      dtype='object')
```

```
# Checking the Data Types
df.dtypes
```

Patient ID	int64
Age	int64
Gender	object
Admission Type	object
Length of Stay	int64
Number of Diagnoses	int64
Blood Pressure	int64
Blood Sugar Levels	int64
Previous Admissions	int64

```
; Readmission          object  
dtype: object
```

Observation

- Categorical columns in our data: Gender, Admission Type, Readmission are the
- The rest of the columns are numerical columns: Patient ID, Age, Length of Stay, Number of Diagnoses, Blood Pressure, Blood Sugar Levels, Previous Admissions

```
# It helps us select categorical columns from our data using its data type  
print("Categorical Columns:", list(df.select_dtypes(include=['object']).columns))
```

```
Categorical Columns: ['Gender', 'Admission Type', 'Readmission']
```

```
# It helps us select numerical columns from our data using its data type  
print("Numerical Columns:", list(df.select_dtypes(include=['int64']).columns))
```

```
Numerical Columns: ['Patient ID', 'Age', 'Length of Stay', 'Number of Diagnoses', 'Blood Pressure', 'Blood Sugar Level  
s', 'Previous Admissions']
```

Step 2: Check the data quality

```
# Counting the number of rows that are duplicated in data  
df.duplicated().sum()
```

```
np.int64(0)
```

```
# Displays more detail of duplications. False means 0 duplications  
df.duplicated()
```

```
0    False  
1    False  
2    False  
3    False  
4    False  
...  
2995   False  
2996   False  
2997   False  
2998   False  
2999   False  
Length: 3000, dtype: bool
```

```
# Checking the number os NULL values in dataset  
df.isnull().sum()
```

```
Patient ID      0  
Age             0  
Gender          0  
Admission Type  0  
Length of Stay   0  
Number of Diagnoses  0  
Blood Pressure    0  
Blood Sugar Levels  0  
Previous Admissions  0  
Readmission      0  
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3000 entries, 0 to 2999  
Data columns (total 10 columns):  
 #   Column           Non-Null Count  Dtype     
 ---  --  
 0   Patient ID      3000 non-null   int64    
 1   Age              3000 non-null   int64    
 2   Gender           3000 non-null   object    
 3   Admission Type  3000 non-null   object    
 4   Length of Stay   3000 non-null   int64    
 5   Number of Diagnoses  3000 non-null   int64    
 6   Blood Pressure   3000 non-null   int64    
 7   Blood Sugar Levels  3000 non-null   int64    
 8   Previous Admissions  3000 non-null   int64    
 9   Readmission      3000 non-null   object    
 dtypes: int64(7), object(3)  
 memory usage: 234.5+ KB
```

```
# Visual inspection for checking impossible values
display(df)
```

Table

	Patient ID	Age	Gender	Admission Type	Length of Stay	Number of Diagnoses
1	1	62	Female	Elective	4	
2	2	65	Male	Emergency	19	
3	3	82	Female	Emergency	18	
4	4	85	Male	Emergency	2	
5	5	85	Female	Elective	19	
6	6	27	Male	Emergency	18	
7	7	39	Male	Elective	21	
8	8	54	Male	Emergency	3	
9	9	88	Female	Elective	6	
10	10	30	Male	Elective	6	
11	11	76	Female	Emergency	13	
12	12	83	Male	Emergency	13	
13	13	57	Female	Emergency	22	
14	14	64	Male	Elective	3	
15						

3,000 rows

Step 3: Describe Each Column

Describe our numerical columns

```
df.describe()
```

	Patient ID	Age	Length of Stay	Number of Diagnoses	Blood Pressure	Blood Sugar Levels	Previous Admissions
count	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000
mean	1500.500000	53.453000	15.226333	4.956667	119.058000	133.536000	1.958333
std	866.169729	20.892996	8.264858	2.575073	34.089265	37.626311	1.401637
min	1.000000	18.000000	1.000000	1.000000	60.000000	51.000000	0.000000
25%	750.750000	35.000000	8.000000	3.000000	89.750000	101.000000	1.000000
50%	1500.500000	53.000000	15.000000	5.000000	119.000000	133.000000	2.000000
75%	2250.250000	72.000000	22.000000	7.000000	147.250000	167.000000	3.000000
max	3000.000000	89.000000	29.000000	9.000000	179.000000	199.000000	4.000000

Describe Categorical Columns

```
# Checking the unique values in the 'Readmission', 'Gender' and 'Admission Type' columns
print(df['Readmission'].unique())
print(df['Gender'].unique())
print(df['Admission Type'].unique())

['No' 'Yes']
['Female' 'Male']
['Elective' 'Emergency']

# Gives us the number(count) of the number of different values in the 'Readmission' 'Gender' and 'Admission Type' col
print(df['Readmission'].nunique())
print(df['Gender'].nunique())
print(df['Admission Type'].nunique())

2
2
2
```

```
# Shows you the unique values in the 'Readmission' column and the number of times they appear
df['Readmission'].value_counts()

Readmission
No      2134
Yes     866
Name: count, dtype: int64
```

```
# Shows you the unique values in the 'Gender' column and the number of times they appear
df['Gender'].value_counts()
```

```
Gender
Male      1555
Female    1445
Name: count, dtype: int64
```

```
# Shows you the unique values in the 'Admission Type' column and the number of times they appear
df['Admission Type'].value_counts()
```

```
Admission Type
Elective   1563
Emergency 1437
Name: count, dtype: int64
```

Step 4: Visualize Single Variables

Univariate Analysis

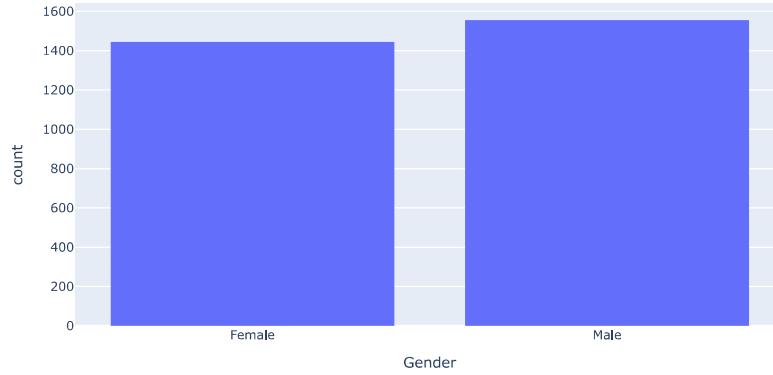
```
# Import plotting libraries
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
df.head(3)
```

	Patient ID	Age	Gender	Admission Type	Length of Stay	Number of Diagnoses	Blood Pressure	Blood Sugar Levels	Previous Admissions	Readmission
0	1	62	Female	Elective	4	5	110	130	1	No
1	2	65	Male	Emergency	19	2	157	81	4	No
2	3	82	Female	Emergency	18	4	74	84	0	No

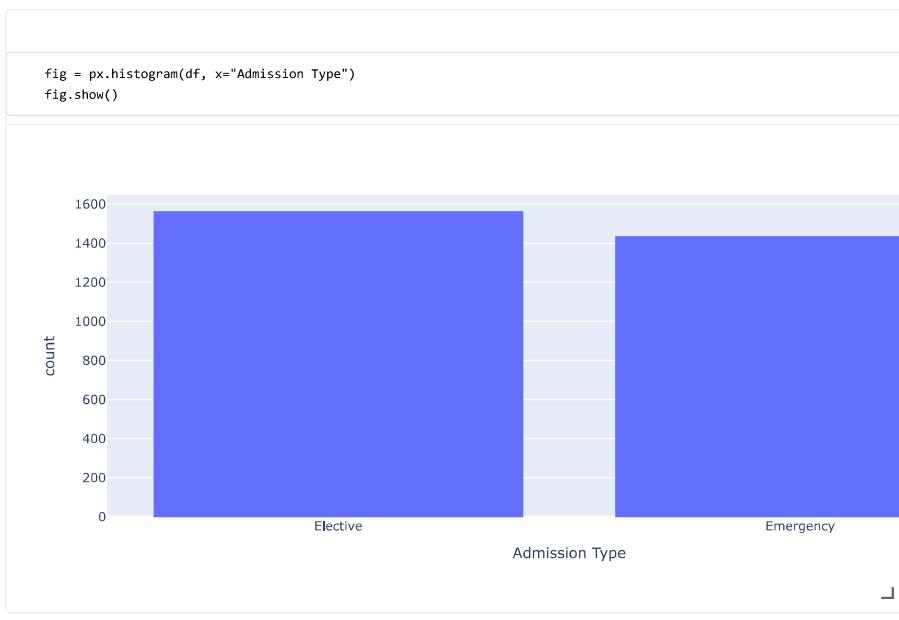
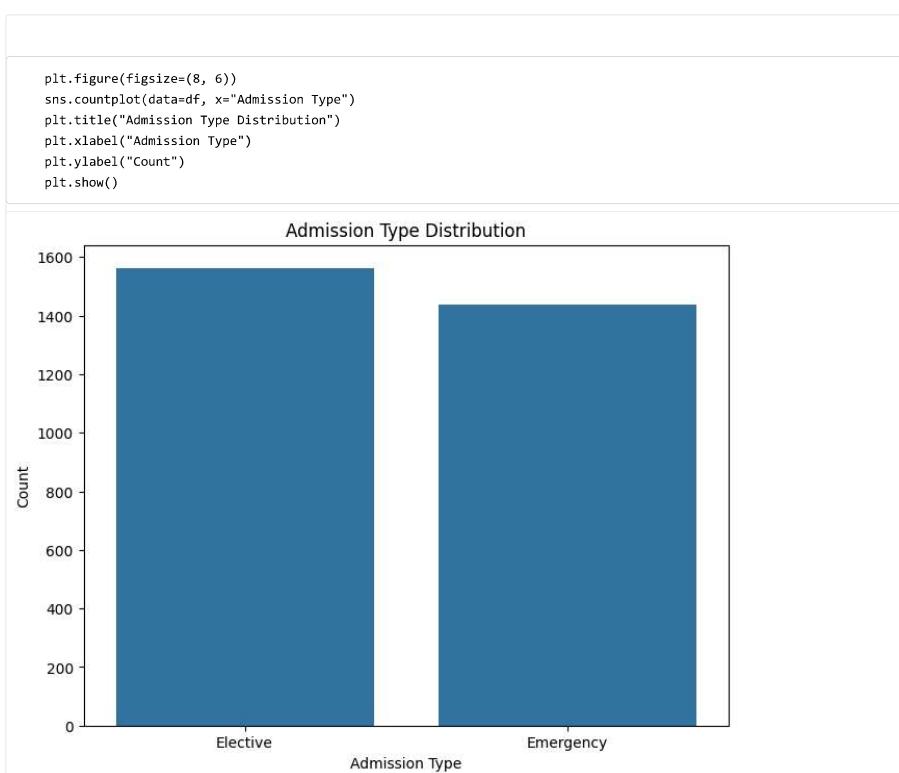
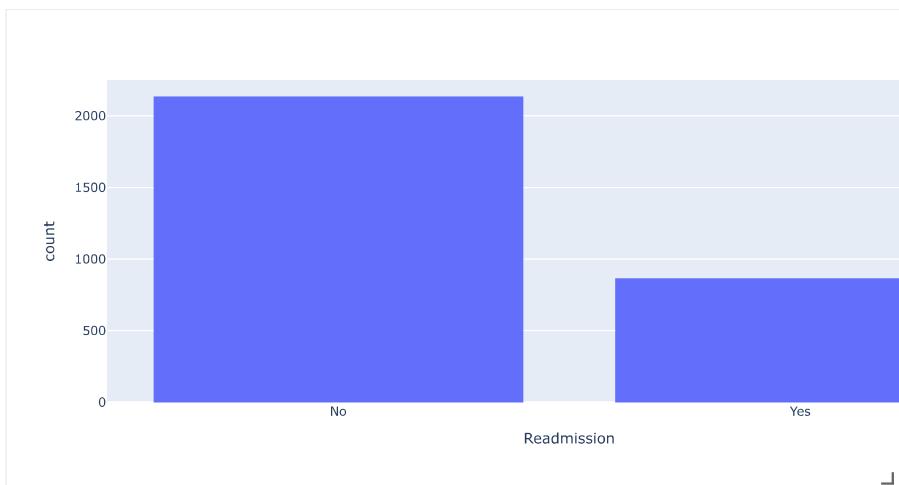
34

```
fig = px.histogram(df, x="Gender")
fig.show()
```



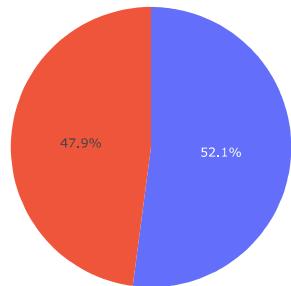
35

```
fig = px.histogram(df, x="Readmission")
fig.show()
```



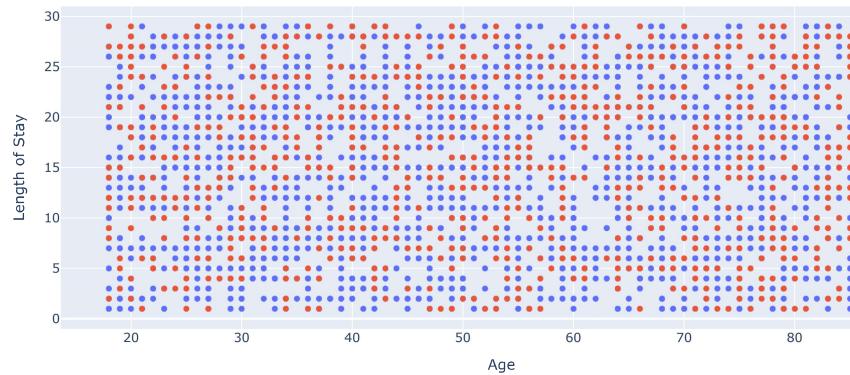
```
fig = px.pie(df, names="Admission Type", title="Admission Type Distribution")
fig.show()
```

Admission Type Distribution



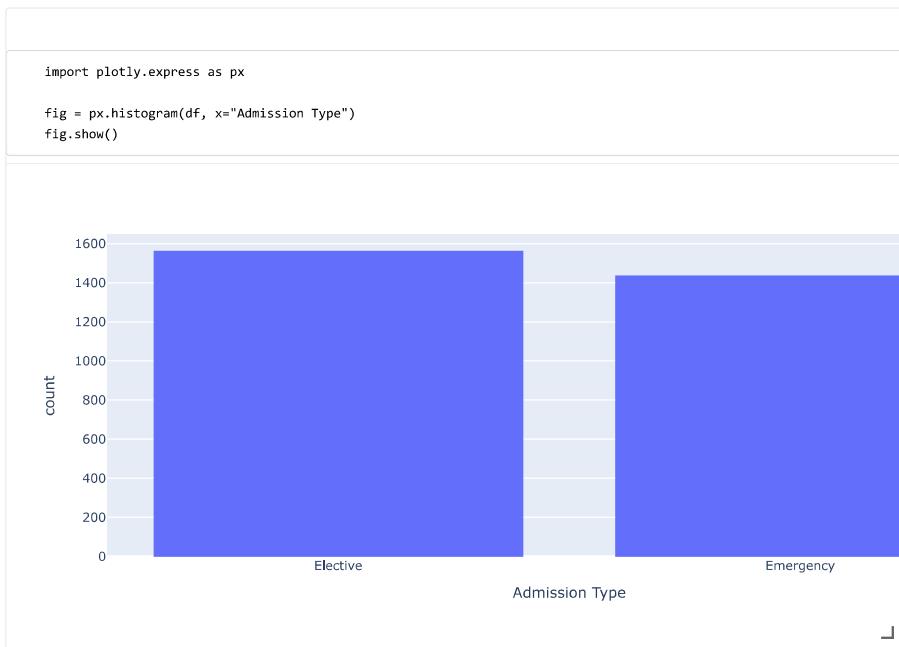
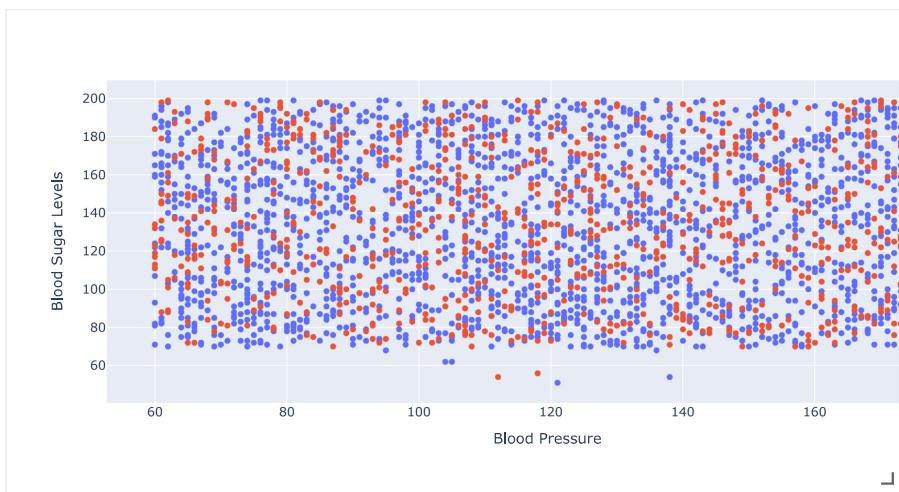
```
import plotly.express as px

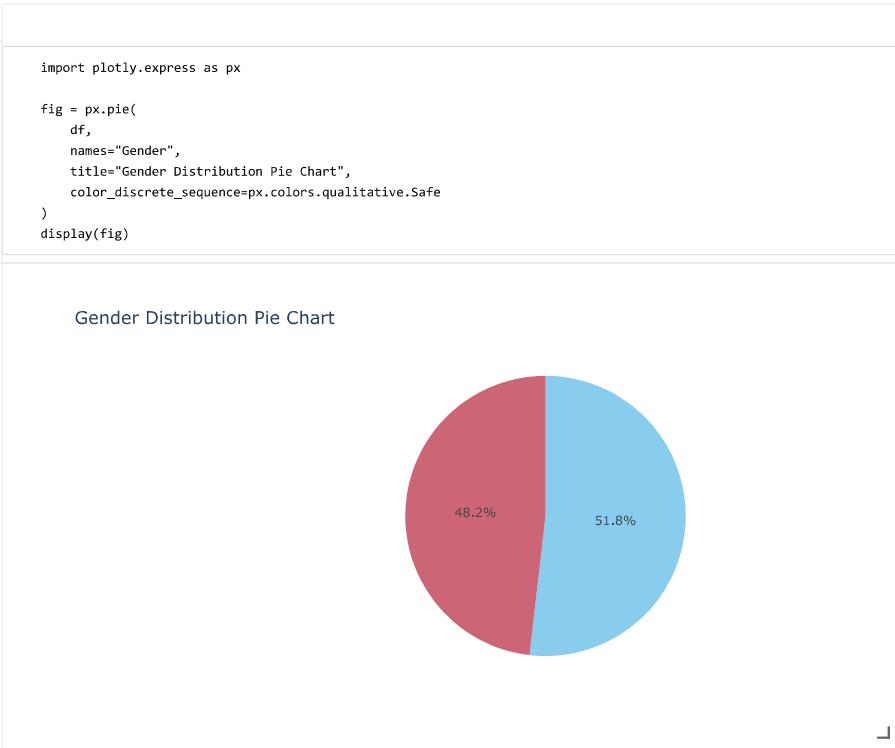
fig = px.scatter(
    df, x = 'Age',
    y = 'Length of Stay',
    color = 'Readmission')
fig.show()
```



```
import plotly.express as px

fig = px.scatter(
    df, x = 'Blood Pressure',
    y = 'Blood Sugar Levels',
    color = 'Readmission',
    hover_name = 'Patient ID')
fig.show()
```





Step 5: Explore Relationships Between Variables

(Bivariate/Multivariate Analysis)

df.head(3)											
Patient ID	Age	Gender	Admission Type	Length of Stay	Number of Diagnoses	Blood Pressure	Blood Sugar Levels	Previous Admissions	Readmission		
0	1	62	Female	Elective	4	5	110	130	1	No	
1	2	65	Male	Emergency	19	2	157	81	4	No	
2	3	82	Female	Emergency	18	4	74	84	0	No	

```

import plotly.figure_factory as ff

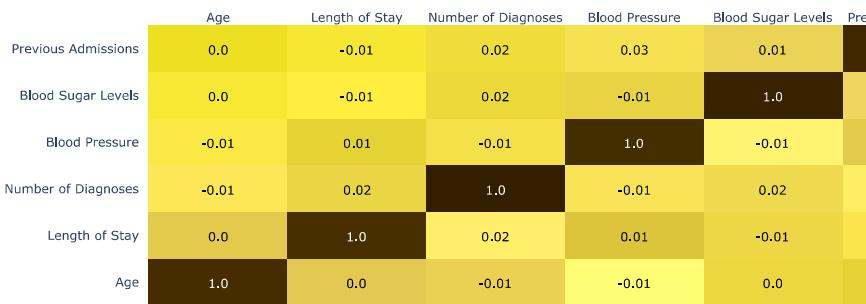
cols = [
    "Age",
    "Length of Stay",
    "Number of Diagnoses",
    "Blood Pressure",
    "Blood Sugar Levels",
    "Previous Admissions"
]

corr_matrix = df[cols].corr().round(2)
fig = ff.create_annotated_heatmap(
    z=corr_matrix.values,
    x=cols,
    y=cols,
    annotation_text=corr_matrix.values.astype(str),
    colorscale='Viridis',
    showscale=True,
    reversescale=True
)
fig.update_layout(title="Correlation Matrix of Selected Features")
fig.show()

```

>  corr_matrix: pandas.core.frame.DataFrame = [Age: float64, Length of Stay: float64 ... 4 more fields]

Correlation Matrix of Selected Features



Observation:

The correlation matrix above displays the pairwise relationships between our selected numerical features: Age, Length of Stay, Number of Diagnoses, Blood Pressure, Blood Sugar Levels, and Previous Admissions. Each cell in the matrix shows the correlation coefficient between two features, ranging from -1 (perfect negative correlation) to +1 (perfect positive correlation), with 0 indicating no linear relationship.

Interpreting the matrix:

- High positive values (close to +1) indicate that as one feature increases, the other tends to increase as well.
- High negative values (close to -1) suggest that as one feature increases, the other tends to decrease.
- Values near 0 imply little to no linear relationship between the features.

This analysis helps identify which features are strongly related and may provide insights for further modeling or feature selection.

```

import plotly.express as px
fig = px.scatter(
    df,
    x="Blood Pressure",
    y="Blood Sugar Levels",
    color="Readmission",
    hover_name="Patient ID",
    size_max=60,

```

)

fig.show()

Feature Engineering

Encoding our categorical columns

```
from sklearn.preprocessing import OneHotEncoder
import pandas as pd

target_col = 'Readmission'
y = df[target_col]
X_raw = df.drop(columns=[target_col])

categorical_cols = X_raw.select_dtypes(include='object').columns
numeric_cols = X_raw.columns.difference(categorical_cols)

ohe = OneHotEncoder(
    handle_unknown='ignore',
    sparse_output=False # Ensures output is a dense array
)

encoded_array = ohe.fit_transform(X_raw[categorical_cols])
encoded_feature_names = ohe.get_feature_names_out(categorical_cols)

encoded_df = pd.DataFrame(
    encoded_array,
    columns=encoded_feature_names,
    index=df.index
)

numeric_df = X_raw[numeric_cols]

X = pd.concat([numeric_df, encoded_df], axis=1)

display(X.head())
display(y.head())
```

> `encoded_df: pandas.core.frame.DataFrame = [Gender_Female: float64, Gender_Male: float64 ... 2 more fields]`
> `numeric_df: pandas.core.frame.DataFrame = [Age: int64, Blood Pressure: int64 ... 5 more fields]`
> `X: pandas.core.frame.DataFrame = [Age: int64, Blood Pressure: int64 ... 9 more fields]`
> `X_raw: pandas.core.frame.DataFrame = [Patient ID: int64, Age: int64 ... 7 more fields]`

Table

0	No
1	No
2	No
3	No
4	No

Name: Readmission, dtype: object

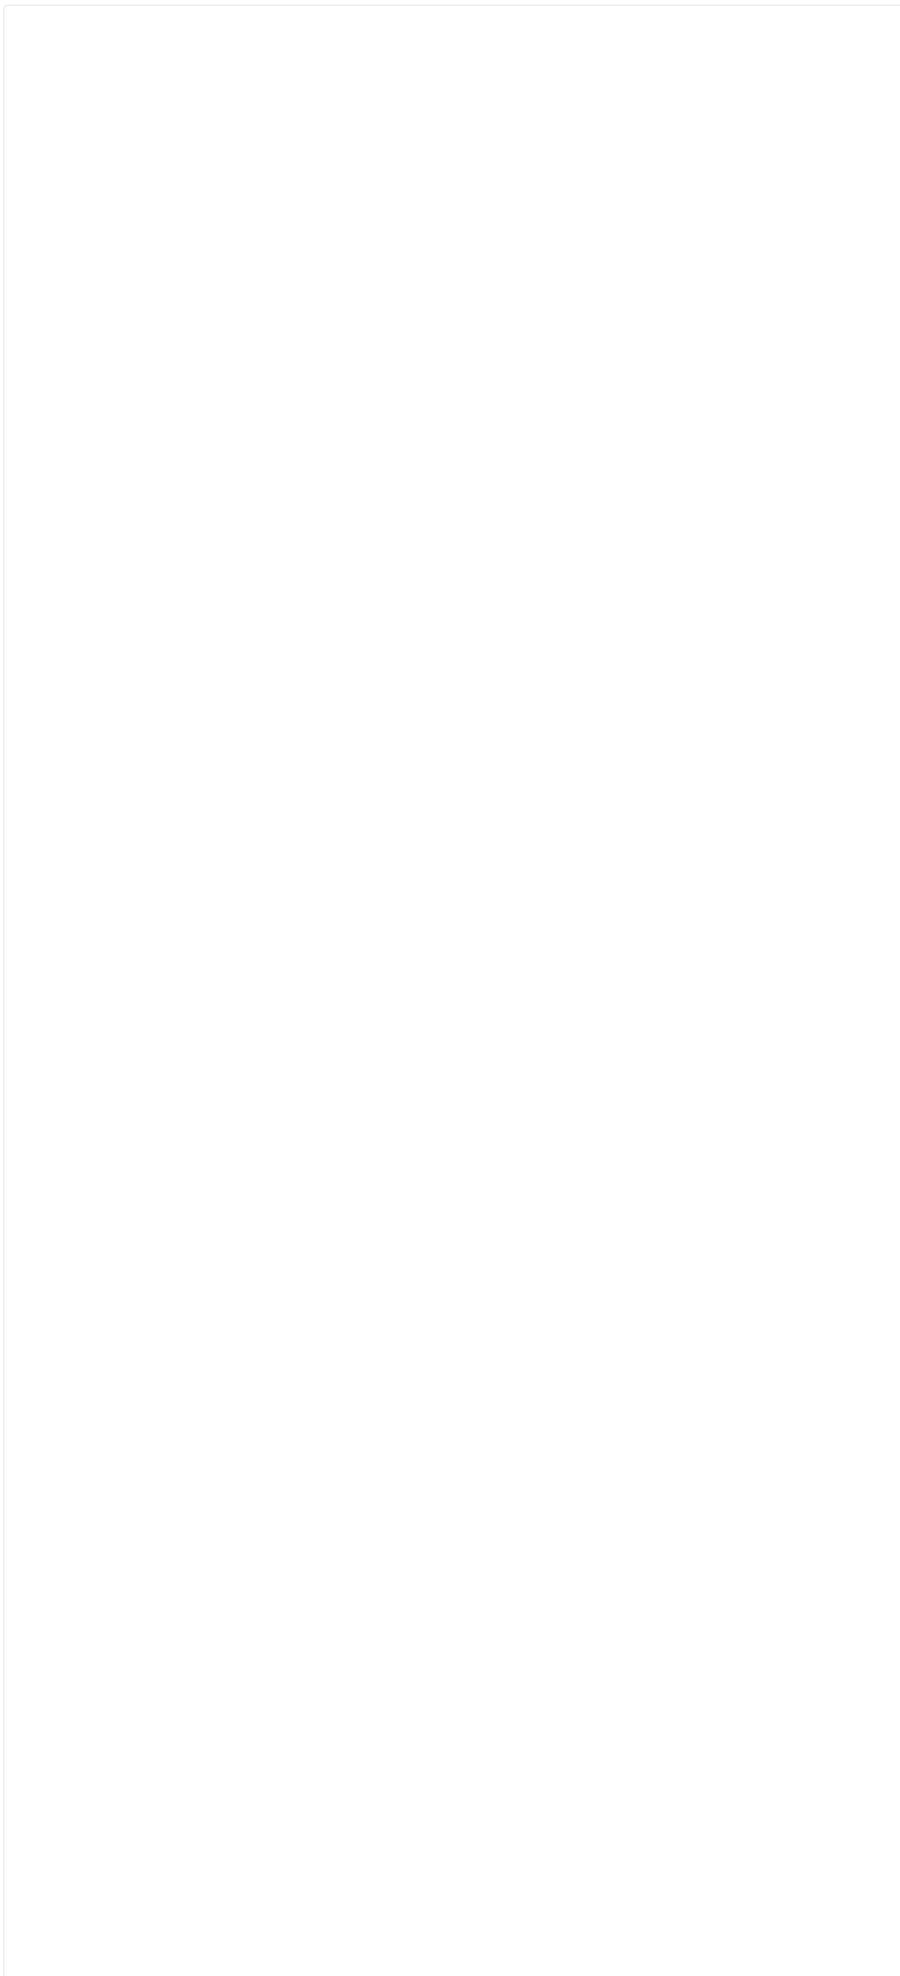
```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,      # 20% test, 80% train
    random_state=42,    # for reproducibility
    stratify=y          # optional but good for classification
)

print(X_train.shape, X_test.shape)
print(y_train.shape, y_test.shape)
```

> `X_test: pandas.core.frame.DataFrame = [Age: int64, Blood Pressure: int64 ... 9 more fields]`
> `X_train: pandas.core.frame.DataFrame = [Age: int64, Blood Pressure: int64 ... 9 more fields]`
(2400, 11) (600, 11)
(2400,) (600,)

Model Training



```

from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    classification_report, confusion_matrix, ConfusionMatrixDisplay
)

# Algorithms
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier

# import matplotlib.pyplot as plt
# import numpy as np

# -----
# Define classification models
# -----
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Random Forest': RandomForestClassifier(random_state=42, n_jobs=-1),
    'Support Vector Classifier': SVC(),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Naive Bayes': GaussianNB(),
    'Gradient Boosting': GradientBoostingClassifier(random_state=42)
}

# -----
# Define evaluation metrics
# -----
metrics = {
    'Accuracy': accuracy_score,
    'Precision': lambda y_true, y_pred: precision_score(
        y_true, y_pred, average='weighted', zero_division=0
    ),
    'Recall': lambda y_true, y_pred: recall_score(
        y_true, y_pred, average='weighted', zero_division=0
    ),
    'F1 Score': lambda y_true, y_pred: f1_score(
        y_true, y_pred, average='weighted', zero_division=0
    ),
}
}

# -----
# Train and evaluate models
# -----
evaluation_results = {}
confusion_matrices = {}

# Get class labels (for nicer axis names in plots)
class_labels = np.unique(y_test)

for name, model in models.items():
    print('=' * 70)
    print(f"Training {name}:\n")

    # Fit model
    model.fit(X_train, y_train)

    # Predict on test set
    y_pred = model.predict(X_test)

    # Store metrics
    evaluation_results[name] = {
        metric_name: metric_func(y_test, y_pred)
        for metric_name, metric_func in metrics.items()
    }

    # Compute and store confusion matrix
    cm = confusion_matrix(y_test, y_pred, labels=class_labels)
    confusion_matrices[name] = cm

    # Print detailed classification report
    print(f"Classification Report for {name}:")
    print(classification_report(y_test, y_pred, zero_division=0))
    print("Confusion Matrix (raw values):")
    print(cm)
    print("-" * 50)

    # -----
    # Plot confusion matrix for this model
    # -----
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_labels)
    disp.plot()
    plt.title(f"Confusion Matrix - {name}")
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    # -----
    # Print summary table

```

```

# -----
print("\nOverall Evaluation Results:")
for name, scores in evaluation_results.items():
    print(f"\n{name}")
    for metric_name, score in scores.items():
        print(f" {metric_name}: {score:.4f}")

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessing.html)

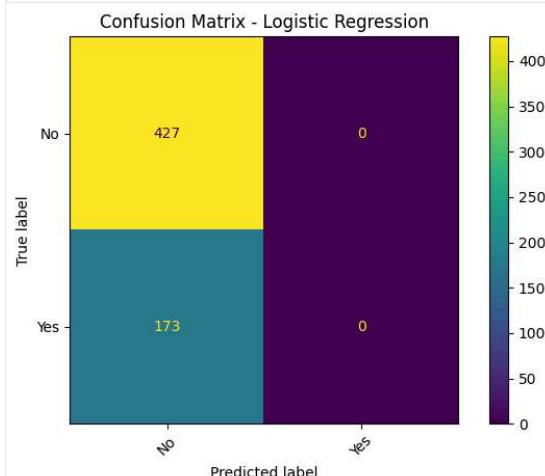
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
No	0.71	1.00	0.83	427
Yes	0.00	0.00	0.00	173
accuracy			0.71	600
macro avg	0.36	0.50	0.42	600
weighted avg	0.51	0.71	0.59	600

Confusion Matrix (raw values):
[[427 0]]

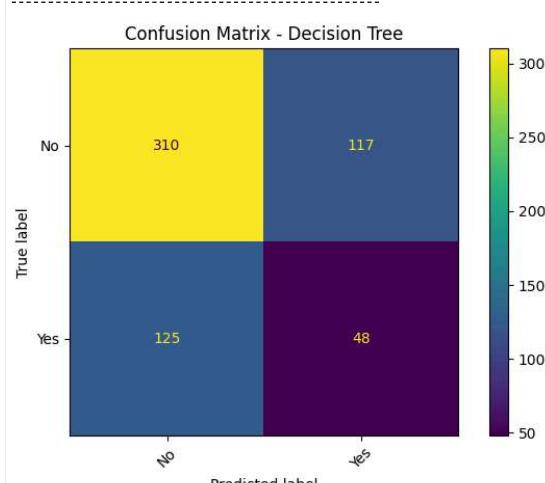


=====

Training Decision Tree:

	precision	recall	f1-score	support
No	0.71	0.73	0.72	427
Yes	0.29	0.28	0.28	173
accuracy			0.60	600
macro avg	0.50	0.50	0.50	600
weighted avg	0.59	0.60	0.59	600

Confusion Matrix (raw values):
[[310 117]
[125 48]]



=====

Training Random Forest:

```

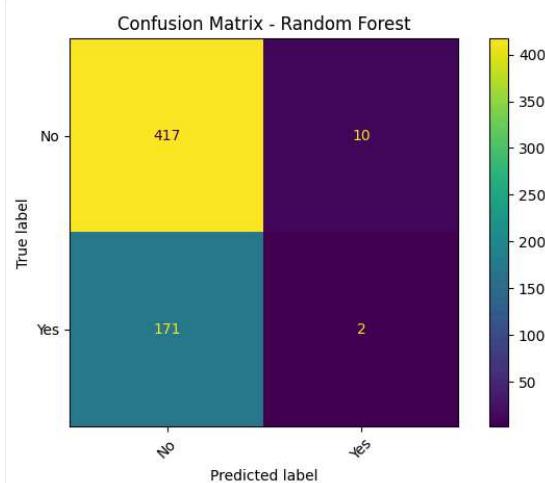
Classification Report for Random Forest:
      precision    recall  f1-score   support

        No       0.71      0.98      0.82     427
       Yes       0.17      0.01      0.02     173

   accuracy                           0.70     600
  macro avg       0.44      0.49      0.42     600
weighted avg       0.55      0.70      0.59     600

```

Confusion Matrix (raw values):
[[417 10]
[171 2]]



Training Support Vector Classifier:

```

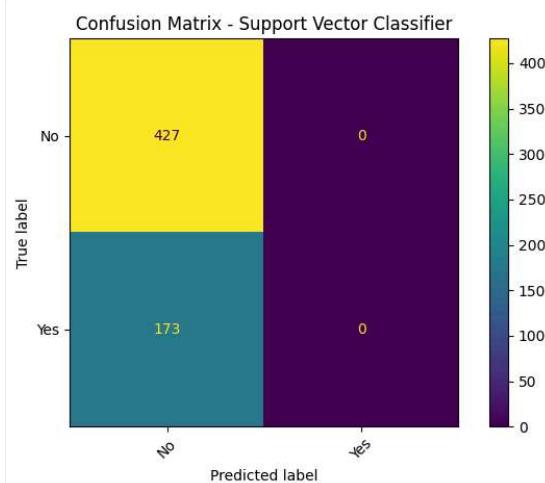
Classification Report for Support Vector Classifier:
      precision    recall  f1-score   support

        No       0.71      1.00      0.83     427
       Yes       0.00      0.00      0.00     173

   accuracy                           0.71     600
  macro avg       0.36      0.50      0.42     600
weighted avg       0.51      0.71      0.59     600

```

Confusion Matrix (raw values):
[[427 0]
[173 0]]



Training K-Nearest Neighbors:

```

Classification Report for K-Nearest Neighbors:
      precision    recall  f1-score   support

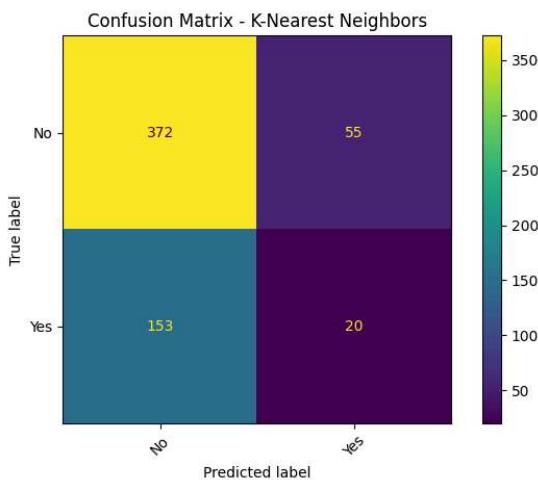
        No       0.71      0.87      0.78     427
       Yes       0.27      0.12      0.16     173

   accuracy                           0.65     600
  macro avg       0.49      0.49      0.47     600
weighted avg       0.58      0.65      0.60     600

```

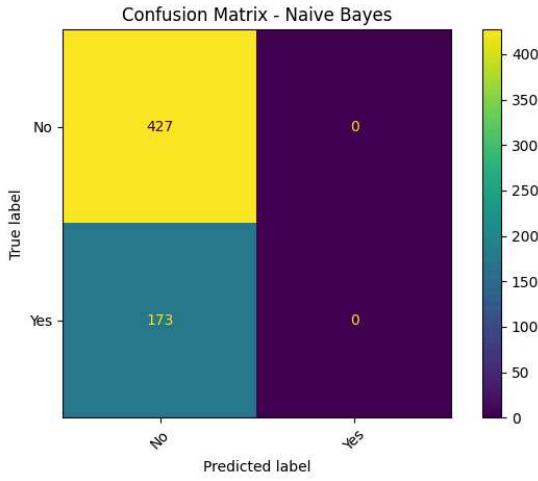
Confusion Matrix (raw values):
[[372 55]

```
[153 20]]
```



```
Classification Report for Naive Bayes:  
precision recall f1-score support  
No 0.71 1.00 0.83 427  
Yes 0.00 0.00 0.00 173  
  
accuracy 0.71 600  
macro avg 0.36 0.50 0.42 600  
weighted avg 0.51 0.71 0.59 600
```

```
Confusion Matrix (raw values):  
[[427 0]  
[173 0]]
```



```
Classification Report for Gradient Boosting:  
precision recall f1-score support  
No 0.71 0.98 0.82 427  
Yes 0.31 0.02 0.04 173  
  
accuracy 0.70 600  
macro avg 0.51 0.50 0.43 600  
weighted avg 0.60 0.70 0.60 600
```

```
Confusion Matrix (raw values):  
[[418 9]  
[169 4]]
```

